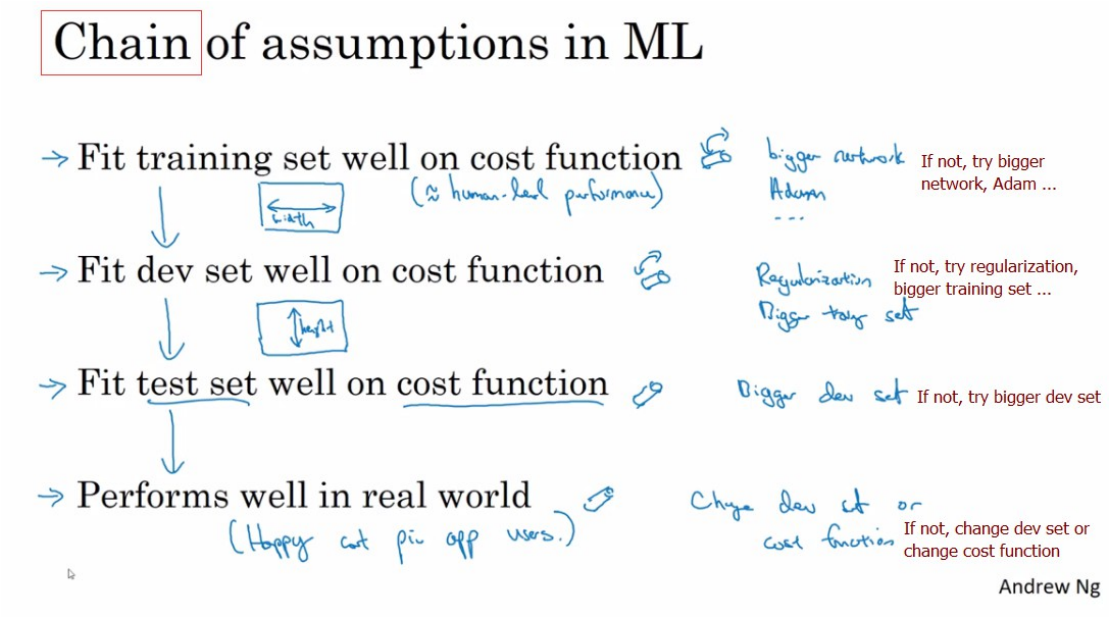

Single real number evaluation metric:

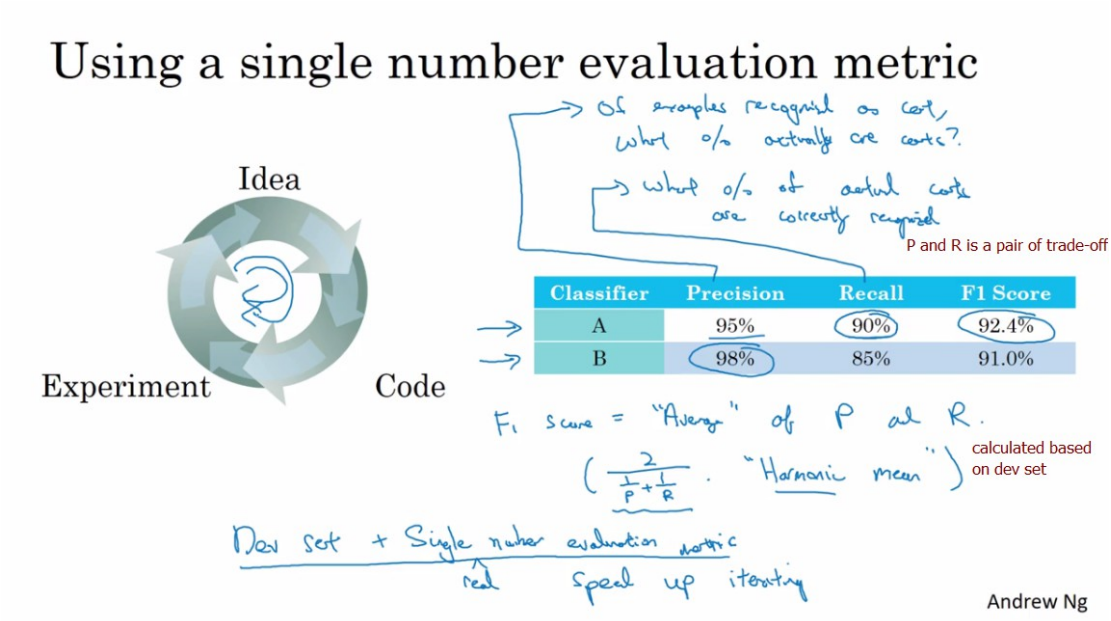

Another example

single real number evaluation metric

Algorithm	US	China	India	Other	Average
A	3%	7%	5%	9%	6%
B	5%	6%	5%	10%	6.5%
C	2%	3%	4%	5%	3.5%
D	5%	8%	7%	2%	5.25%
E	4%	5%	2%	4%	3.75%
F	7%	11%	8%	12%	9.5%

Optimizing and satisfying metric:

Another cat classification example

Classifier	Accuracy	Running time
A	90%	80ms
B	92%	95ms
C	95%	1,500ms

optimizing — satisficing

Cost = accuracy − 0.5 × running Time

try to maximize accuracy

with satisficing to running Time ≤ 100 ms.

N metrics: 1 optimizing, N−1 satisficing

Wakewords / Trigger words

Alexa, OK Google, Hey Siri, 你好百度

accuracy. #false positive

Maximize accuracy. s.t. ≤ 1 false positive every 24 hours.

Train, Dev and Test set splitting:

Guideline

Same distribution

Choose a dev set and test set to reflect data you expect to get in the future and consider important to do well on.

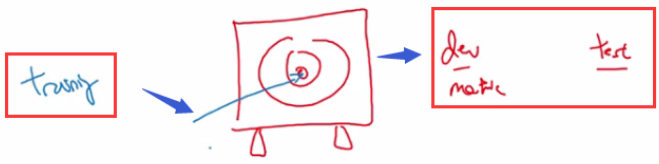

Old way of splitting data

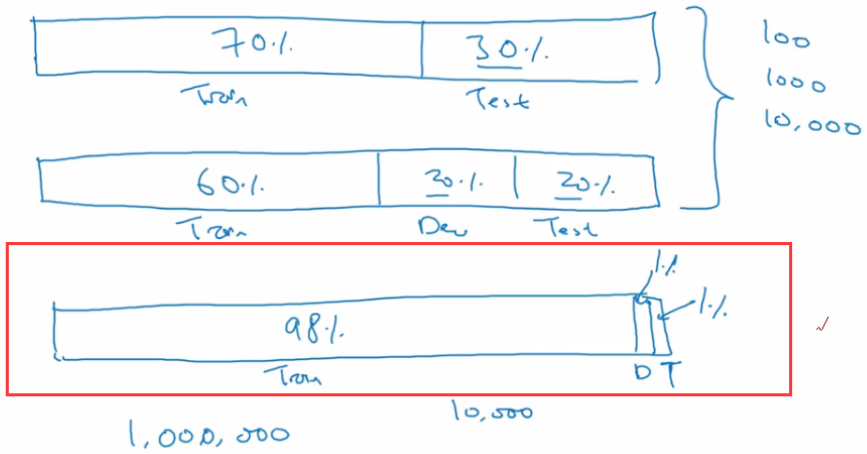

Size of test set

⇒ Set your test set to be big enough to give high confidence in the overall performance of your system.

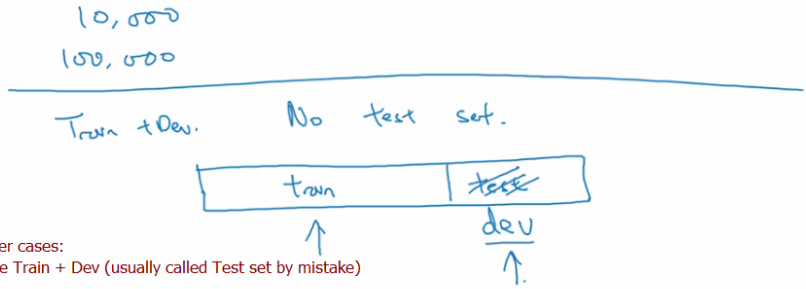

And some other cases:
people just use Train + Dev (usually called Test set by mistake)

When to change Dev/Test set and metric:

Cat dataset examples

Metric + Dev : Prefer A
But You/users : Prefer B.

Metric: classification error

Algorithm A: 3% error → pornographic

Because model A treats porn and non-porn equally

✓ Algorithm B: 5% error

function counts the number of non-match cases

$$\text{Error}: \frac{1}{m_{dev}} \sum_{i=1}^{m_{dev}} \mathbb{1}\{ y_{pred}^{(i)} \neq y^{(i)} \}$$

$y_{pred}^{(i)}$: predicted value (0/1)

that's the sign that we need to change evaluation metric

Cat dataset examples

Metric + Dev : Prefer A
You/users : Prefer B.

→ Metric: classification error

Algorithm A: 3% error → pornographic

✓ Algorithm B: 5% error

$$\text{Error}: \frac{1}{\sum w^{(i)}} \cancel{\frac{1}{m_{dev}}} \sum_{i=1}^{m_{dev}} w^{(i)} \mathbb{1}\{ y_{pred}^{(i)} \neq y^{(i)} \}$$

$y_{pred}^{(i)}$: predicted value (0/1)

$$\rightarrow w^{(i)} = \begin{cases} 1 & \text{if } x^{(i)} \text{ is non-porn} \\ 10 & \text{if } x^{(i)} \text{ is porn} \end{cases}$$

Add more weight to the Error term

Orthogonalization for cat pictures: anti-porn

→ 1. So far we've only discussed how to define a metric to evaluate classifiers. ← Place target

1st step, place the target

→ 2. Worry separately about how to do well on this metric.

↑ Aim/shoot at target

$$\rightarrow J = \frac{1}{\sum w^{(i)}} \sum_{i=1}^{m} w^{(i)} \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

2nd step, define a cost function to shot the target

Another example

Algorithm A: 3% error

Algorithm B: 5% error

→ Dev/test

→ User images

If doing well on your metric + dev/test set does not correspond to doing well on your application, change your metric and/or dev/test set.

Human Level Performance:

Comparing to human-level performance

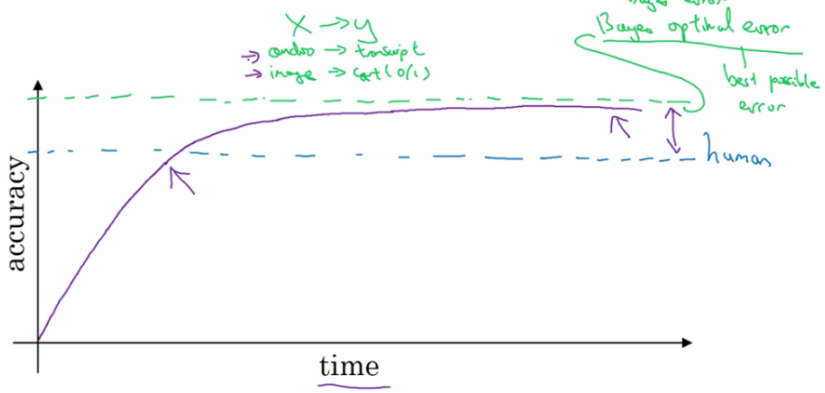

Avoidable Bias:

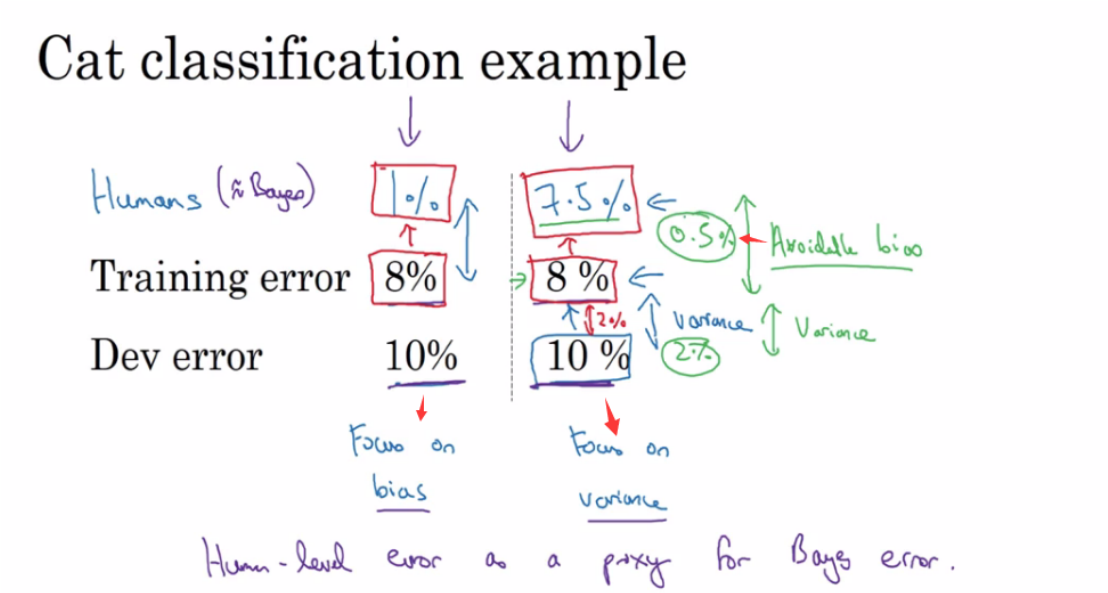

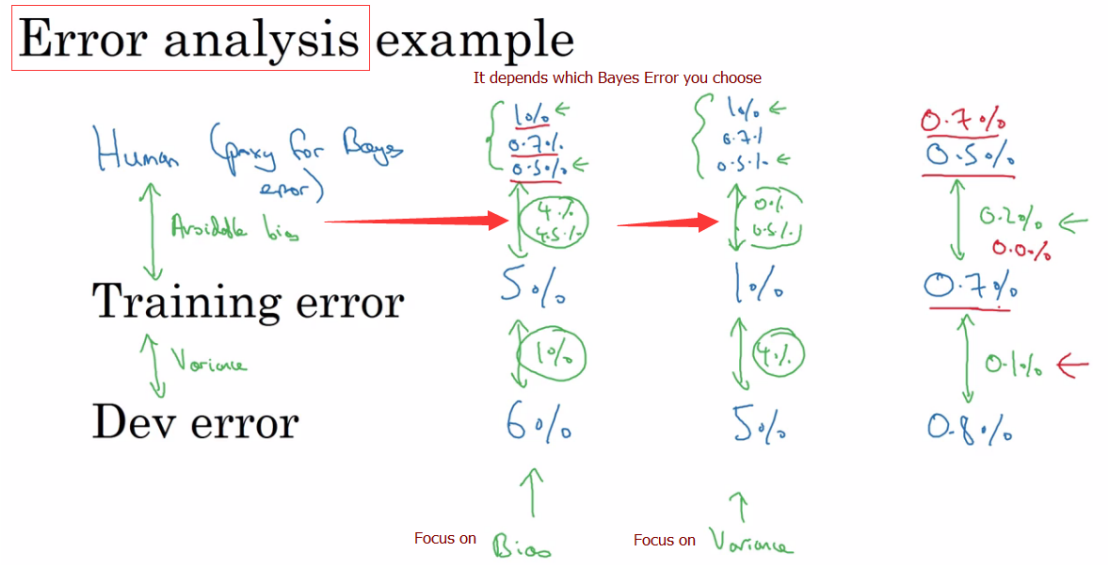

Improving Model Performance:

The two fundamental assumptions of supervised learning

1. You can fit the training set pretty well.

~ Avoidable bias

2. The training set performance generalizes pretty well to the dev/test set.

~ Variance

Reducing (avoidable) bias and variance